

Abstractions and Frameworks for Deep Learning: a Discussion

—

Caffe, Torch, Theano, TensorFlow, et al.

Rémi Emonet

Saintélyon Deep Learning Workshop – 2015-11-26

Disclaimer



Introductory Poll

- Did you ever use?

- Caffe
- Theano
- Lasagne
- Torch
- TensorFlow
- Other?

- Any experience to share?



Overview

- Deep Learning?
- Abstraction in Frameworks
- A Tour of Existing Framework
- More Discussions?



Overview

- Deep Learning?
- Abstraction in Frameworks
- A Tour of Existing Framework
- More Discussions?

Finding Parameters of a Function (supervised)

- Notations
 - Input i
 - Output o
 - Function f given
 - Parameters θ to be learned
 - We suppose: $o = f_{\theta}(i)$
- How to optimize it: how to find the best θ ?
 - need some regularity assumptions
 - usually, at least differentiability
- Remark: a more generic view
 - $o = f_{\theta}(i) = f(\theta, i)$

Gradient Descent

- We want to find the best parameters
 - we suppose: $o = f_{\theta}(i)$
 - we have examples of inputs i^n and target output t^n
 - we want to minimize the sum of errors $L(\theta) = \sum_n L(f_{\theta}(i^n), t^n)$
 - we suppose f and L are differentiable
- Gradient descent (gradient = vector of partial derivatives)
 - start with a random θ^0
 - compute the gradient and update $\theta^{t+1} = \theta^t - \gamma \nabla_{\theta} L(\theta)$
- Variations
 - stochastic gradient descent (SGD)
 - conjugate gradient descent
 - BFGS
 - L-BFGS
 - ...

Finding Parameters of a “Deep” Function

- Idea

- f is a composition of functions
- 2 layers: $o = f_{\theta}(i) = f_{\theta^2}^2(f_{\theta^1}^1(i))$
- 3 layers: $o = f_{\theta}(i) = f_{\theta^3}^3(f_{\theta^2}^2(f_{\theta^1}^1(i)))$
- K layers: $o = f_{\theta}(i) = f_{\theta^K}^K(\dots f_{\theta^3}^3(f_{\theta^2}^2(f_{\theta^1}^1(i)))) \dots$
- with all f_l differentiable

- How can we optimize it?

- The chain rule!

- Many versions (with $F = f \circ g$)

- $(f \circ g)' = (f' \circ g) \cdot g'$
- $F'(x) = f'(g(x))g'(x)$
- $\frac{df}{dx} = \frac{df}{dg} \cdot \frac{dg}{dx}$

Finding Parameters of a “Deep” Function

- Reminders: K layers: $o = f_{\theta}(i) = f_{\theta^K}^K(\dots f_{\theta^3}^3(f_{\theta^2}^2(f_{\theta^1}^1(i))))\dots$

- minimize the sum of errors $L(\theta) = \sum_n L(f_{\theta}(i^n), t^n)$

- chain rule $\frac{df}{dx} = \frac{df}{dg} \cdot \frac{dg}{dx}$

- Goal: compute $\nabla_{\theta} L$ for gradient descent

- $\nabla_{\theta^K} L = \frac{dL}{d\theta^K} = \frac{dL}{df^K} \frac{df^K}{d\theta^K}$

- $\nabla_{\theta^{K-1}} L = \frac{dL}{d\theta^{K-1}} = \frac{dL}{df^K} \frac{df^K}{df^{K-1}} \frac{df^{K-1}}{d\theta^{K-1}}$

- $\nabla_{\theta^1} L = \frac{dL}{d\theta^1} = \frac{dL}{df^K} \frac{df^K}{df^{K-1}} \dots \frac{df^2}{df^1} \frac{df^1}{d\theta^1}$

- $\frac{dL}{df^K}$: gradient of the loss with respect to its input ✓

- $\frac{df^k}{df^{k-1}}$: gradient of a function with respect to its input ✓

- $\frac{df^k}{d\theta^k}$: gradient of a function with respect to its parameters ✓

Deep Learning and Composite Functions

- Deep Learning?
 - NN can be deep, CNN can be deep
 - “any” composition of differentiable function can be optimized with gradient descent
 - some other models are also deep... (hierarchical models, etc)
- Evaluating a composition $f_{\theta}(i) = f_{\theta^K}^K(\dots f_{\theta^3}^3(f_{\theta^2}^2(f_{\theta^1}^1(i)))\dots)$
 - “forward pass”
 - evaluate successively each function
- Computing the gradient $\nabla_{\theta}L$ (for gradient descent)
 - compute the input (\$o\$) gradient (from the output error)
 - for each f_1, f_2, \dots
 - compute the parameter gradient (from the output gradient)
 - compute the input gradient (from the output gradient)

Back to “seeing parameters as inputs”

- Parameters (θ^k)
- Just another input of f_k
- Can be rewritten, e.g. as $f_k(\theta_k, x)$
- More generic
 - inputs can be constant
 - inputs can be parameters
 - inputs can be produced by another function (e.g. $f(g(x), h(x))$)



Overview

- Deep Learning?
- **Abstraction in Frameworks**
- A Tour of Existing Framework
- More Discussions?

Function/Operator/Layer

- The functions that we can use for f_k
- Many choices
 - fully connected layers
 - convolutions layers
 - activation functions (element-wise)
 - soft-max
 - pooling
 - ...
- Loss Functions: same with no parameters
- In the wild
 - Torch module
 - Theano operator

Data/Blob/Tensor

- The data: input, intermediate result, parameters, gradient, ...
- Usually a tensor (n-dimensional matrices)
- In the wild
 - Torch tensor
 - Theano tensor, scalars, numpy arrays

Overview

- Deep Learning?
- Abstraction in Frameworks
- A Tour of Existing Framework
- More Discussions?

Contenders

- Caffe
- Torch
- Theano
- Lasagne
- Tensor Flow
- Deeplearning4j
- ...

Overview

- Basics
 - install CUDA/Cublas/OpenBlas
 - blob/tensors, blocks/layers/loss, parameters
 - cuDNN
 - open source
- Control flow
 - define a composite function (graph)
 - choice of an optimizer
 - forward, backward
- Extend
 - write a new operator/module
 - "forward"
 - "backward": gradParam, gradInput

- "made with expression, speed, and modularity in mind"
- "developed by the Berkeley Vision and Learning Center (BVLC)"
- "released under the BSD 2-Clause license"
- C++
- layers-oriented <http://caffe.berkeleyvision.org/tutorial/layers.html>
- plaintext protocol buffer schema (prototxt) to describe models (and so save them too)
- 1,068 / 7,184 / 4,077

Torch7

- By
 - Ronan Collobert (Idiap, now Facebook)
 - Clement Farabet (NYU, now Madbits now Twitter)
 - Koray Kavukcuoglu (Google DeepMind)
- Lua (+ C)
 - need to learn
 - easy to embed
- Layer-oriented
 - easy to use
 - difficult to extend, sometimes (merging sources)
- 418 / 3,267 / 757

Theano

- "is a Python library"
- "allows you to define, optimize, and evaluate mathematical expressions"
- "involving multi-dimensional arrays"
- "efficient symbolic differentiation"
- "transparent use of a GPU"
- "dynamic C code generation"
- Use symbolic expressions: **reasoning on the graph**
 - write numpy-like code
 - no forced "layered" architecture
 - computation graph
- 263 / 2,447 / 878

Lasagne (Keras, etc)

- Overlay to Theano
- Provide layer API close to caffe/torch etc
- Layer-oriented
- 133 / 1,401 / 342

Tensor Flow

- By Google, Nov. 2015
- Selling points
 - easy to move from a cluster to a mobile phone
 - easy to distribute
- Currently slow?
- Not fully open yet?
- 1,303 / 13,232 / 3,375

Deeplearning4j

- “Deep Learning for Java, Scala & Clojure on Hadoop, Spark & GPUs”
- Apache 2.0-licensed
- Java
- High level (layer-oriented)
- Typed API
- 236 / 1,648 / 548

Overview

- Deep Learning?
- Abstraction in Frameworks
- A Tour of Existing Framework
- More Discussions?

Be creative!
anything differentiable
can be tried!

How to choose a framework?

Any experience to share?



Thanks

Abstractions and Frameworks
for Deep Learning:
a Discussion

–

Caffe, Torch, Theano, TensorFlow, et al.

Rémi Emonet

Saintélyon Deep Learning Workshop – 2015-11-26